

Class 8: Machine Learning Mini Project

Danika

In today's mini-project we will explore a complete analysis using unsupervised learning techniques covered in class (clustering and PCA for now).

The data itself comes from the Wisconsin Breast Cancer Diagnosis Data Set FNA breast biopsy data.

Save your input data file into your Project directory

```
fna.data <- "WisconsinCancer.csv"
```

Complete the following code to input the data and store as wisc.df

```
wisc.df <- read.csv(fna.data, row.names=1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430

843786	0.12780	0.17000	0.1578	0.08089
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se
842302	0.2419	0.07871	1.0950	0.9053
842517	0.1812	0.05667	0.5435	0.7339
84300903	0.2069	0.05999	0.7456	0.7869
84348301	0.2597	0.09744	0.4956	1.1560
84358402	0.1809	0.05883	0.7572	0.7813
843786	0.2087	0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se
842302	153.40	0.006399	0.04904	0.05373
842517	74.08	0.005225	0.01308	0.01860
84300903	94.03	0.006150	0.04006	0.03832
84348301	27.23	0.009110	0.07458	0.05661
84358402	94.44	0.011490	0.02461	0.05688
843786	27.19	0.007510	0.03345	0.03672
	concave.points_se			
842302				0.01587
842517				0.01340
84300903				0.02058
84348301				0.01867
84358402				0.01885
843786				0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302	0.11890			
842517	0.08902			
84300903	0.08758			
84348301	0.17300			
84358402	0.07678			
843786	0.12440			

Remove the diagnosis column and keep it in a separate vector for later # We can use -1 here to remove the first column

```
diagnosis <- as.factor(wisc.df[,1])
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780

	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean
842302	0.27760	0.3001	0.14710	0.2419
842517	0.07864	0.0869	0.07017	0.1812
84300903	0.15990	0.1974	0.12790	0.2069
84348301	0.28390	0.2414	0.10520	0.2597
84358402	0.13280	0.1980	0.10430	0.1809
843786	0.17000	0.1578	0.08089	0.2087

	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302	0.07871	1.0950	0.9053	8.589	153.40
842517	0.05667	0.5435	0.7339	3.398	74.08
84300903	0.05999	0.7456	0.7869	4.585	94.03
84348301	0.09744	0.4956	1.1560	3.445	27.23
84358402	0.05883	0.7572	0.7813	5.438	94.44
843786	0.07613	0.3345	0.8902	2.217	27.19

	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	0.006399	0.04904	0.05373	0.01587
842517	0.005225	0.01308	0.01860	0.01340
84300903	0.006150	0.04006	0.03832	0.02058
84348301	0.009110	0.07458	0.05661	0.01867
84358402	0.011490	0.02461	0.05688	0.01885
843786	0.007510	0.03345	0.03672	0.01137

	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75

	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249

	concavity_worst	concave.points_worst	symmetry_worst
842302	0.7119	0.2654	0.4601
842517	0.2416	0.1860	0.2750
84300903	0.4504	0.2430	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4000	0.1625	0.2364
843786	0.5355	0.1741	0.3985

	fractal_dimension_worst
842302	0.11890
842517	0.08902
84300903	0.08758
84348301	0.17300
84358402	0.07678
843786	0.12440

Exploratory data analysis

The first step of any data analysis, unsupervised or supervised, is to familiarize yourself with the data.

Q1. How many observations (patients) are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

Q3. How many variables/features in the data are suffixed with `_mean`?

First find the column names

```
colnames(wisc.data)
```

```
[1] "radius_mean"           "texture_mean"
[3] "perimeter_mean"       "area_mean"
[5] "smoothness_mean"      "compactness_mean"
[7] "concavity_mean"       "concave.points_mean"
[9] "symmetry_mean"        "fractal_dimension_mean"
[11] "radius_se"            "texture_se"
[13] "perimeter_se"         "area_se"
[15] "smoothness_se"        "compactness_se"
[17] "concavity_se"         "concave.points_se"
[19] "symmetry_se"          "fractal_dimension_se"
[21] "radius_worst"         "texture_worst"
[23] "perimeter_worst"      "area_worst"
[25] "smoothness_worst"     "compactness_worst"
[27] "concavity_worst"      "concave.points_worst"
[29] "symmetry_worst"       "fractal_dimension_worst"
```

Next I need to search within the column names for the “`_mean`” pattern. The ‘`grep()`’ function might help here.

```
inds <- grep("_mean", colnames((wisc.data)))
length(inds)
```

```
[1] 10
```

Q. How many dimensions are in this dataset?

```
ncol(wisc.data)
```

```
[1] 30
```

Principal Component Analysis

First do we need to scale the data before PCA or not.

Check column means and standard deviations

```
round(apply(wisc.data,2,sd),3)
```

radius_mean	texture_mean	perimeter_mean
3.524	4.301	24.299
area_mean	smoothness_mean	compactness_mean
351.914	0.014	0.053
concavity_mean	concave.points_mean	symmetry_mean
0.080	0.039	0.027
fractal_dimension_mean	radius_se	texture_se
0.007	0.277	0.552
perimeter_se	area_se	smoothness_se
2.022	45.491	0.003
compactness_se	concavity_se	concave.points_se
0.018	0.030	0.006
symmetry_se	fractal_dimension_se	radius_worst
0.008	0.003	4.833
texture_worst	perimeter_worst	area_worst
6.146	33.603	569.357
smoothness_worst	compactness_worst	concavity_worst
0.023	0.157	0.209
concave.points_worst	symmetry_worst	fractal_dimension_worst
0.066	0.062	0.018

Looks like we need to scale.

Perform PCA on wisc.data by completing the following code

```
wisc.pr <- prcomp( wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14

Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27%

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PCs capture 72%

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs capture 91%

Interpreting PCA results


Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

There is a distinct difference between 2 sets of data. We need to better organize the data to make sense of it.

PC Plot


We need to make out plot of PC1 vs PC2 (aka score plot, PC-plot, etc.) The main results of PCA...

```
biplot(wisc.pr)
```



Scatter plot observations by components 1 and 2

```
plot( wisc.pr$x[,1],wisc.pr$x[,2], col = diagnosis ,
      xlab = "PC1", ylab = "PC2")
```

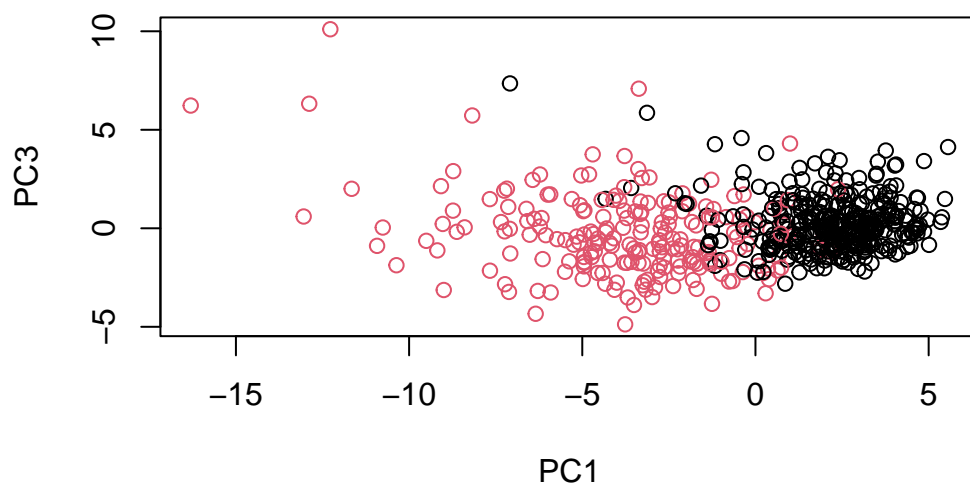



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

The red and the black localization of points are around the same areas on the graph. The majority of the red dots is on the left of the majority of the black dots.

Repeat for components 1 and 3


```
plot(wisc.pr$x[,1], wisc.pr$x[,3], col = diagnosis,  
     xlab = "PC1", ylab = "PC3")
```



```
library(ggplot2)

pc <- as.data.frame(wisc.pr$x)
pc$diagnosis <- diagnosis

ggplot(pc)+
  aes(PC1, PC2, col=diagnosis)+
  geom_point()
```



Variance Explained

Calculate the variance of each principal component by squaring the sdev component of `wisc.pr` (i.e. `wisc.pr$sdev^2`). Save the result as an object called `pr.var`.

Calculate variance of each component

```
pr.var <- wisc.pr$sdev^2  
head(pr.var)
```

```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```


Calculate the variance explained by each principal component by dividing by the total variance explained of all principal components. Assign this to a variable called `pve` and create a plot of variance explained for each principal component.

Variance explained by each principal component: pve

```
pve <- pr.var / sum(pr.var)
```


Plot variance explained for each principal component

```
plot(pve, xlab = "Principal Component",  
     ylab = "Proportion of Variance Explained",  
     ylim = c(0, 1), type = "o")
```



Alternative scree plot of the same data, note data driven y-axis

```
barplot(pve, ylab = "Precent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

How much do the original variables contribute to the new PCs that we have calculated? To get at this data we can look at the ‘`$rotation`’ component of the returned PCA object

```
head(wisc.pr$rotation[,1:3])
```

	PC1	PC2	PC3
radius_mean	-0.2189024	0.23385713	-0.008531243
texture_mean	-0.1037246	0.05970609	0.064549903
perimeter_mean	-0.2275373	0.21518136	-0.009314220
area_mean	-0.2209950	0.23107671	0.028699526
smoothness_mean	-0.1425897	-0.18611302	-0.104291904
compactness_mean	-0.2392854	-0.15189161	-0.074091571

Focus in on PC1


```
wisc.pr$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

There is complicated mix of variables that go together to make up PC1 i.e. there are many of the original variables that together contribute highly to PC1.

```
loadings <- as.data.frame(wisc.pr$rotation)

ggplot(loadings)+
  aes(PC1, rownames(loadings))+
  geom_col()
```



Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

5 PCs for 85%

Hierarchical clustering

The goal of this section is to do hierarchical clustering of the original data.


First we will scale the data, then distance matrix, then hclust

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(dist(scale(wisc.data)))
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

Height 19

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```



```
dist(scale(wisc.data))
hclust (*, "complete")
```

Cut this tree to yield cluster membership vector with ‘`cutree()`’ function. How well do the two clusters separate the M and B diagnoses?

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)
table(wisc.hclust.clusters, diagnosis)
```

diagnosis

```
wisc.hclust.clusters  B  M
      1 12 165
      2  2  5
      3 343 40
      4  0  2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

I think a better cluster vs diagnoses match would be at 10 since there is a bigger distinction between benign and malignant data points.


```
x <- 10
wisc.hclust.clusters <- cutree(wisc.hclust, k=x)
table(wisc.hclust.clusters, diagnosis)
```

```
              diagnosis
wisc.hclust.clusters  B  M
      1 12 86
      2  0 59
      3  0  3
      4 331 39
      5  0 20
      6  2  0
      7 12  0
      8  0  2
      9  0  2
     10  0  1
```

Try clustering in 3 PCs, that is PC1, PC2, and PC3 as input

```
d <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(d, method="ward.D2")
plot(wisc.pr.hclust)
```



Cluster Dendrogram



```
hclust (*, "ward.D2")
```

```
d <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(d, method="complete")
plot(wisc.pr.hclust)
```


Cluster Dendrogram



d
hclust (*, "complete")

```
d <- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(d, method="single")  
plot(wisc.pr.hclust)
```


Cluster Dendrogram



d
hclust (*, "single")

```
d <- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(d, method="average")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram



d
hclust (*, "average")

Q13. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning

My favorite method is the 'complete' because it distributes the branches more evenly and is easier to see the different branches.

Q15. How well does the newly created model with four clusters separate out the two diagnoses?


The newly created model with the four cluster separated out the points visually with the plot in different colors, but the table does not make a clear distinction between the patients with malignant and benign conditions.

#Combine methods: PCA and HCLUST My PCA results were interesting as they showed a separation of M and B samples along PC1.

Let's cut this tree into four groups/clusters


```
d <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(d, method="ward.D2")
plot(wisc.pr.hclust)
```

Cluster Dendrogram




```
grps2 <- cutree(wisc.pr.hclust, k=4)
```

```
plot(wisc.pr$x[,1],wisc.pr$x[,2], col=diagnosis)
```



I want to cluster my PCA results - that is use 'wisc.pr\$x' as input 'hclust()'


```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=grps2)
```



And my tree result figure

```
plot(wisc.pr.hclust)
```

Cluster Dendrogram



```
d
hclust (*, "ward.D2")
```

Compare to actual diagnoses

```
table(grps2, diagnosis)
```

```
      diagnosis
grps2  B    M
1      0  111
2     24   68
3    184   32
4    149    1
```

```
(179+333)/nrow(wisc.data)
```

```
[1] 0.8998243
```